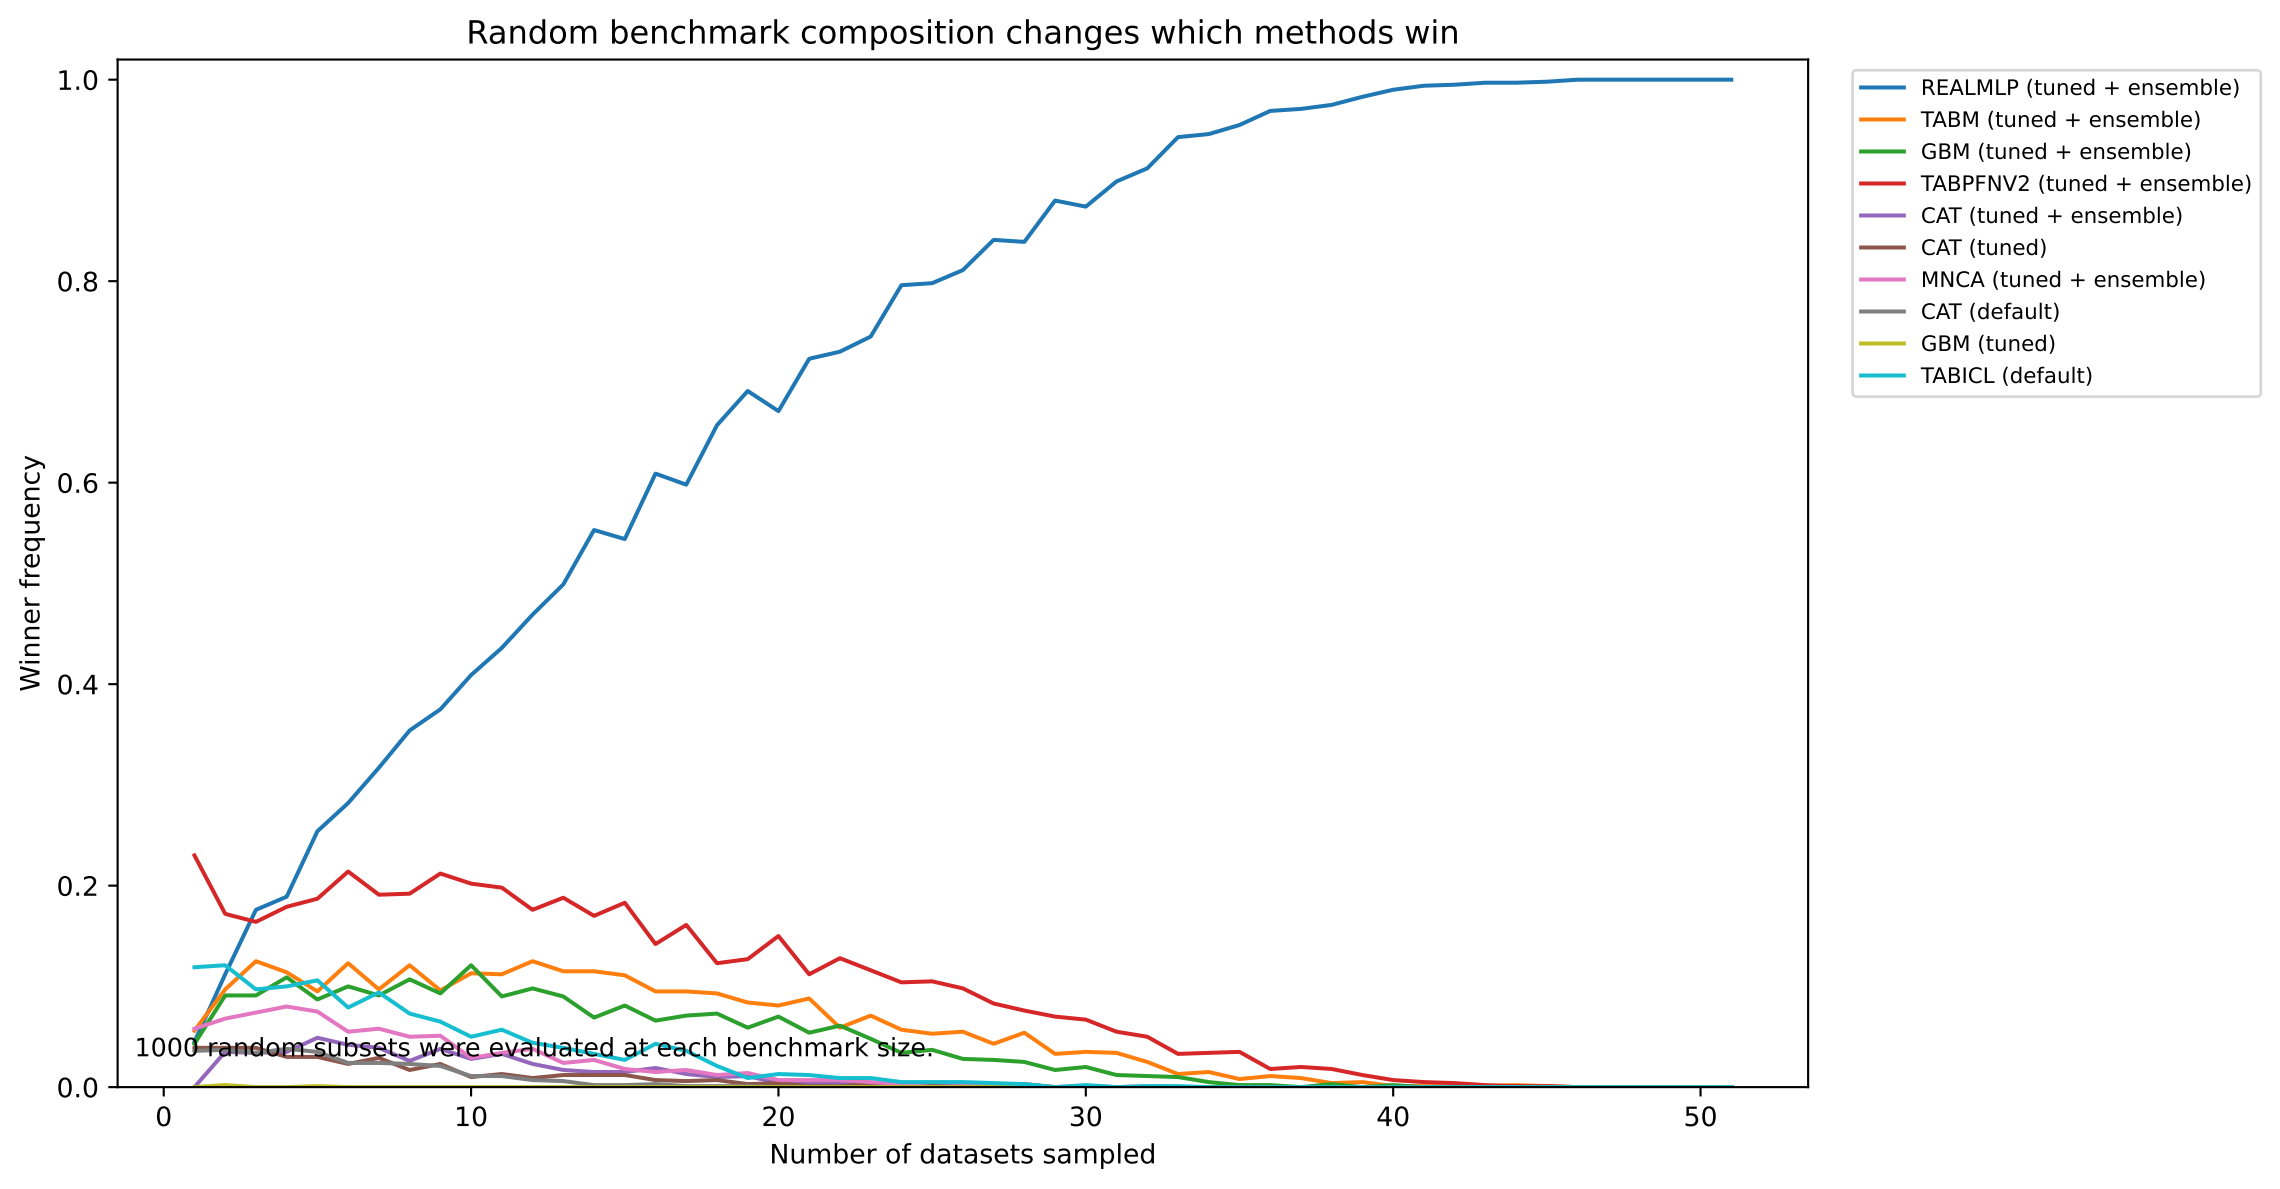
Random benchmark composition changes which methods win
Winner frequency
Number of datasets sampled
1000 random subsets were evaluated at each benchmark size.
REALMLP (tuned + ensemble)
TABM (tuned + ensemble)
GBM (tuned + ensemble)
TABPFNV2 (tuned + ensemble)
CAT (tuned + ensemble)
CAT (tuned)
MNCA (tuned + ensemble)
CAT (default)
GBM (tuned)
TABICL (default)